

# Implementation of an Iterative Linear Quadratic Regulator (iLQR)

Gabriel Desfrene   Antoine Groudiev

January 14, 2025



# Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

Conclusion

# Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

Conclusion

## General formulation

- Dynamics function:

$$x_{t+1} = f(x_t, u_t)$$

- Goal: minimize a quadratic cost function
- Cost function:

$$J(u) = \sum_{t=0}^{T-1} \left( x_t^\top Q x_t + u_t^\top R u_t \right) + \frac{1}{2} (x_T - x^*)^\top Q_f (x_T - x^*)$$


- $Q$ : state cost matrix
- $Q_f$ : final state cost matrix
- $R$ : control cost matrix

## Example: Simple Pendulum

- State:  $x = [\theta \ \dot{\theta}]$
- Control:  $u$ , torque applied to the pendulum
- Dynamics: physical laws (simulator)
- Target:  $x = [0 \ 0]$
- Cost function:

$$J(u) = \frac{1}{2} \left( \theta_f^2 + \dot{\theta}_f^2 \right) + \frac{1}{2} \int_0^T r u^2(t) dt$$

corresponding to  $Q_f = I_2$ ,  $Q = 0_2$ ,  $R = rI_1$




## Example: Cartpole

- State:  $x = [y \ \theta \ \dot{y} \ \dot{\theta}]$
- Control:  $u$ , force applied to the cart
- Dynamics: physical laws (simulator)
- Target:  $x = [0 \ 0 \ 0 \ 0]$
- Cost function:

$$J(u) = \frac{1}{2} \left( \theta_f^2 + \dot{\theta}_f^2 + y_f^2 + \dot{y}_f^2 \right) + \frac{1}{2} \int_0^T r u^2(t) dt$$

corresponding to  $Q_f = I_4$ ,  $Q = 0_4$ ,  $R = rI_1$



# Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

Conclusion

## General idea

- iLQR is an iterative algorithm
- Start with an initial trajectory
- Iteratively improve it using a local linear approximation
- Stop when the trajectory converges



## Linearizing the dynamics

The equation  $x_{t+1} = f(x_t, u_t)$  is linearized (at each step) as:

$$\delta x_{t+1} = A_t \delta x_t + B_t \delta u_t$$

with:

- $A_t$ : Jacobian of  $f$  with respect to  $x$  evaluated at  $(x_t, u_t)$
- $B_t$ : Jacobian of  $f$  with respect to  $u$  evaluated at  $(x_t, u_t)$

We are in LQR (Linear Quadratic Regulator, cf. TP5) setup!

## Trajectory refinement using LQR

1. **Forward pass:** compute the successive states  $(x_t)$  for the current controls  $(u_t)$ , and the corresponding cost  $J$
2. **Backward pass:** compute the gains, i.e. how much we should change the controls in each direction to minimize the cost
3. **Forward rollout:** apply the gains to the controls to obtain a new trajectory
4. Repeat until convergence

# Computing the Jacobians

## Finite differences method

We want to compute:

- $A_t = \frac{\partial f}{\partial x}(x_t, u_t)$ , i.e. how much the state at time  $t + 1$  changes when we slightly change the state at time  $t$
- $B_t = \frac{\partial f}{\partial u}(x_t, u_t)$ , i.e. how much the state at time  $t + 1$  changes when we slightly change the control at time  $t$

In a black box setting, we can use finite differences:

$$[A_t]_i \approx \frac{f(x_t + \varepsilon e_i, u_t) - f(x_t - \varepsilon e_i, u_t)}{2\varepsilon}$$
$$[B_t]_i \approx \frac{f(x_t, u_t + \varepsilon e_i) - f(x_t, u_t - \varepsilon e_i)}{2\varepsilon}$$

for some small  $\varepsilon$  and the canonical basis  $(e_i)$

## Computing the Jacobians

Using Pinocchio

# Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

Conclusion

## What language to use?

# Python

- Easy to use
- Bindings for many libraries
- Embarrassingly slow

# C++

- Fast
- Not very funny

## Rust

- Fast
- Very funny



## API Basic usage



## Plan

## Problem statement

## The iLQR algorithm

## Our implementation

## Demonstration time

## Conclusion

# Demonstration time!

## Plan

## Problem statement

## The iLQR algorithm

## Our implementation

## Demonstration time

## Conclusion

Problem statement  
○○○○

The iLQR algorithm  
○○○○○○

Our implementation  
○○○○

Demonstration time  
○○

Conclusion  
○●