

Implementation of an Iterative Linear Quadratic Regulator (iLQR)

Gabriel Desfrene Antoine Groudiev

January 15, 2025



Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

General formulation

- Dynamics function:

$$x_{t+1} = f(x_t, u_t)$$

- Goal: minimize a quadratic cost function
- Cost function:

$$J(u) = \sum_{t=0}^{T-1} \left(x_t^\top Q x_t + u_t^\top R u_t \right) + \frac{1}{2} (x_T - x^*)^\top Q_f (x_T - x^*)$$


- Q : state cost matrix
- Q_f : final state cost matrix
- R : control cost matrix

Example: Simple Pendulum

- State: $x = [\theta \ \dot{\theta}]$
- Control: u , torque applied to the pendulum
- Dynamics: physical laws (simulator)
- Target: $x = [0 \ 0]$
- Cost function:

$$J(u) = \frac{1}{2} \left(\theta_f^2 + \dot{\theta}_f^2 \right) + \frac{1}{2} \int_0^T r u^2(t) dt$$

corresponding to $Q_f = I_2$, $Q = 0_2$, $R = rI_1$




Example: Cartpole

- State: $x = [y \ \theta \ \dot{y} \ \dot{\theta}]$
- Control: u , force applied to the cart
- Dynamics: physical laws (simulator)
- Target: $x = [0 \ 0 \ 0 \ 0]$
- Cost function:

$$J(u) = \frac{1}{2} \left(\theta_f^2 + \dot{\theta}_f^2 + y_f^2 + \dot{y}_f^2 \right) + \frac{1}{2} \int_0^T r u^2(t) dt$$

corresponding to $Q_f = I_4$, $Q = 0_4$, $R = rI_1$



Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

General idea

- iLQR is an iterative algorithm
- Start with an initial trajectory
- Iteratively improve it using a local linear approximation
- Stop when the trajectory converges

Linearizing the dynamics

The equation $x_{t+1} = f(x_t, u_t)$ is linearized (at each step) as:

$$\delta x_{t+1} = A_t \delta x_t + B_t \delta u_t$$

with:

- A_t : Jacobian of f with respect to x evaluated at (x_t, u_t)
- B_t : Jacobian of f with respect to u evaluated at (x_t, u_t)

We are in LQR (Linear Quadratic Regulator, cf. TP5) setup!

Trajectory refinement using LQR

1. **Forward pass:** compute the successive states (x_t) for the current controls (u_t), and the corresponding cost J
2. **Backward pass:** compute the gains, i.e. how much we should change the controls in each direction to minimize the cost
3. **Forward rollout:** apply the gains to the controls to obtain a new trajectory
4. Repeat until convergence

For the complete derivations, see [1] or [3].

Computing the Jacobians

Finite differences method

We want to compute:

- $A_t = \frac{\partial f}{\partial x}(x_t, u_t)$, i.e. how much the state at time $t + 1$ changes when we slightly change the state at time t
- $B_t = \frac{\partial f}{\partial u}(x_t, u_t)$, i.e. how much the state at time $t + 1$ changes when we slightly change the control at time t

In a black box setting, we can use finite differences:

$$[A_t]_i \approx \frac{f(x_t + \varepsilon e_i, u_t) - f(x_t - \varepsilon e_i, u_t)}{2\varepsilon}$$
$$[B_t]_i \approx \frac{f(x_t, u_t + \varepsilon e_i) - f(x_t, u_t - \varepsilon e_i)}{2\varepsilon}$$

for some small ε and the canonical basis (e_i)

Computing the Jacobians

Using Pinocchio

```
# Jacobians of the Articulated-Body algorithm
```

```
J1_q, J1_v, J1_u = pin.computeABADerivatives(model, data_sim, q, v, u)
```

```
# compute the Jacobians of the integration on SE(...)
```

```
J2_q, J2_v_2 = pin.dIntegrate(model, q, v_2 * dt)
```

Tricks for practical convergence

- **Gradient clipping:** limit the size of the control updates norm to α to avoid divergence

$$\delta u_t = \frac{\delta u_i}{\max\left(1, \frac{\|\delta u_i\|}{\alpha}\right)}$$

- **Gaussian initialization:** start with a small random control sequence instead of a zero sequence

$$u_t \sim \mathcal{N}(0, \Sigma)$$

- **Regularization based on Levenberg-Marquardt:** When inverting the Q_{uu} matrix, add a dynamic regularization term $\mu > 0$ to ensure positive definiteness

$$Q_{uu} = U\Sigma U^T \rightarrow Q_{uu}^{-1} = U\Sigma' U^T \quad \text{with} \quad \Sigma' = \begin{cases} 0 & \text{if } \sigma_i \leq 0 \\ \frac{1}{\sigma_i + \mu} & \text{otherwise} \end{cases}$$

Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

What language to use?

Python

- Easy to use
- Support for many libraries
- Embarrassingly slow

C++

- "Fast"
- Not very funny

Rust

- "Fast"
- Easy bindings for Python
- Very funny

Therefore, we chose to have a Rust core with Python bindings

From Rust to Python, and the other way around

- Instantiate the solver in Python
- Use Python libraries to define the dynamics
- The Rust solver does the computations, and calls the Python `dynamics` function and the Pinocchio functions for the Jacobians
- Supports both methods for computing the Jacobians

API Basic usage

```
def dynamics(x, u):  
    return ... # simulator  
  
Q = np.zeros((state_dim, state_dim)) # state cost  
Qf = np.eye(state_dim) # final state cost  
R = 1e-5 * np.eye(control_dim) # control cost (minimize the energy)  
  
s = ilqr.ILQRSolver(state_dim, control_dim, Q, Qf, R)  
target = np.zeros(state_dim) # upright pendulum with no velocity  
output = s.solve(np.concatenate((q0, v0)), target, dynamics, time_steps=N,  
                 gradient_clip=10.0, # max norm of the gradient  
                 initialization=0.5) # std of the Gaussian initialization
```

Plan

Problem statement

The iLQR algorithm

Our implementation

Demonstration time

Demonstration time!

References

- [1] Brian Jackson and Taylor Howell. *iLQR Tutorial*. Sept. 2019. URL: https://rexlab.ri.cmu.edu/papers/iLQR_Tutorial.pdf.
- [2] Weiwei Li and Emanuel Todorov. “Iterative linear quadratic regulator design for nonlinear biological movement systems”. In: *First International Conference on Informatics in Control, Automation and Robotics*. Vol. 2. SciTePress. 2004, pp. 222–229.
- [3] Harley Wiltzer. *iLQR Without Obfuscation*. Feb. 2020. URL: <https://harwiltz.github.io/posts/20200201-ilqr/index.html>.